

Near-Optimal Space (MGA) Exploration in GreenBubble

Implementation summary — what we built, what we reused, and what we would like your feedback on

GreenBubble project (DTU)
prepared for a discussion with Aleksander Grochowicz

August 2026

Abstract

GreenBubble is a PyPSA-based capacity-expansion model of an industrial Power-to-X hub (electrolysis, methanolisation, biogas upgrading, storage, symbiosis). This note summarises how we implemented *Modelling to Generate Alternatives* (MGA) / near-optimal space exploration on top of it, following Neumann & Brown (2021) for the near-optimal space itself and Grochowicz et al. (2023) for the multi-year robustness extension. It is written for a working discussion: Section 4 is deliberately explicit about *exactly* what we reused from PyPSA's own MGA machinery versus what we implemented from scratch, since we did not have access to a Grochowicz-authored reference implementation and would value comparing notes.

Contents

1	Motivation	2
2	Mathematical background	2
2.1	The cost-optimal LP and the ε -near-optimal space	2
2.2	Sign-robust budget for price-driven (negative-objective) runs	2
2.3	Projection onto technology-capacity dimensions	3
2.4	Three tiers	3
2.5	Shared cost bound and the Chebyshev centre	3
3	Implementation architecture	3
3.1	Dimension auto-derivation	3
3.2	The constraint-correctness problem	4
3.3	Stochastic-network support	5
4	What we reused versus what we wrote from scratch	5
4.1	Reused directly via PyPSA's public API (not copied, called)	5
4.2	Implemented from scratch (GreenBubble-original)	5
4.3	Verification against the reference implementation	6
4.4	Reuse plan: <code>near-optimal_dev3</code>	6
4.5	The honest summary	7
5	Validation status	7
6	Example results (Tutorial 5)	7

7 Discussion points for the meeting

11

1 Motivation

A cost-optimal capacity-expansion solve returns exactly one design. In practice, many *other* designs are within a few percent of that cost and can differ substantially in technology mix. Exploring this *near-optimal feasible space* answers questions the single optimum cannot: which technologies are load-bearing versus fully substitutable, how much room exists to steer the design toward secondary goals, and whether one design stays good across several weather/price years rather than being over-fitted to the one year it was solved for.

GreenBubble needed this for an additional reason beyond the usual single-country capacity-expansion setting: it is optimised in two distinct *modes* — **demand** mode (minimise cost to meet fixed product demands, objective ≥ 0) and **price** mode (maximise profit selling at fixed prices, which shows up in the LP as *minimising* a cost objective that is typically *negative*, since revenue exceeds cost). Any near-optimal formulation had to be correct in both regimes, and also compose with GreenBubble’s stochastic (multi-scenario) optimisation mode. None of this changes the underlying mathematics of Neumann & Brown or Grochowicz et al. — it just meant we could not use the reference formulation completely as-is and had to generalise the sign convention (see Section 2.2).

2 Mathematical background

2.1 The cost-optimal LP and the ε -near-optimal space

Write the capacity-expansion model as the linear program

$$\min_x c^\top x \quad \text{s.t.} \quad Ax \leq b, \quad (1)$$

where x collects every decision variable (investment *and* operation), $c^\top x$ is total system cost, and c_{opt} is its optimal value. Following Neumann & Brown (2021) and Grochowicz et al. (2023), the **ε -near-optimal feasible space** is the slice of the feasible region whose cost is within a relative slack ε of the optimum:

$$X_\varepsilon = \{x : Ax \leq b, \quad c^\top x \leq (1 + \varepsilon) c_{\text{opt}}\}. \quad (2)$$

2.2 Sign-robust budget for price-driven (negative-objective) runs

GreenBubble adds exactly this extra constraint to the model (the *budget*), but in a sign-robust form so it also holds when $c_{\text{opt}} < 0$ (price mode, net revenue):

$$c^\top x \leq c_{\text{opt}} + \varepsilon |c_{\text{opt}}| = \begin{cases} (1 + \varepsilon) c_{\text{opt}}, & c_{\text{opt}} \geq 0 \quad (\text{demand mode}) \\ (1 - \varepsilon) c_{\text{opt}}, & c_{\text{opt}} < 0 \quad (\text{price mode}) \end{cases} \quad (3)$$

Both branches keep c_{opt} itself inside the feasible band and loosen it by the same relative amount ε . Here c_{opt} is the *optimised LP objective*, **n.objective** — not a re-derived capex+opex statistic. For a **stochastic** network (PyPSA’s multi-scenario mode) it is automatically the probability-weighted *expected* cost $\sum_s p_s c_s^\top x_s$ already, so Eq. (3) defines the near-optimal space consistently for demand, price, *and* stochastic runs without any separate scenario-aware bookkeeping.

2.3 Projection onto technology-capacity dimensions

Exploring the full-dimensional X_ε is neither tractable nor useful for interpretation. We project onto a handful of *technology capacity* coordinates $y = \sigma(x)$ — each y_i is the summed installed capacity of one technology (the σ -aggregation map of Grochowicz et al. §2.2). Exploring a single direction d in this reduced space means re-solving the original model once with the ordinary cost objective replaced by $\pm d^\top \sigma(x)$, subject to the budget of Eq. (3) — this is exactly *Modelling to Generate Alternatives*.

2.4 Three tiers

Tier 1 — ranges. For each dimension, minimise and maximise it alone under the budget: the per-technology min/max installed capacity. $\min > 0 \Rightarrow \textit{must-have}$; $\max \approx 0 \Rightarrow \textit{must-avoid}$.

Tier 2 — hull. Sample many search directions d in the selected capacity space, solve each, and take the convex hull of the resulting extreme points — the approximated near-optimal polytope.

Tier 3 — robustness. Solve the cost-optimum and explore the hull independently for several weather/price years, intersect the resulting per-year polytopes, and take the **Chebyshev centre** (deepest interior point) of the intersection as the design that stays near-optimal *simultaneously* across all years (Grochowicz et al. §2.3–2.4).

2.5 Shared cost bound and the Chebyshev centre

For Tier 3, Grochowicz et al. (eq. 9) note that intersecting per-year spaces that were each bounded by their *own* optimum $c_{\text{opt}}(i)$ is not directly comparable across years with different absolute cost levels; a single shared bound

$$c^* = \max_i c_{\text{opt}}(i) \quad (4)$$

can instead be applied to every year’s budget, so all years are relaxed relative to the *same* anchor. GreenBubble exposes this as `robustness.cost_bound: max` (default) versus `per_year`.

Given m per-year polytopes (each a `scipy.spatial.ConvexHull`), their facet inequalities $a_i^\top y \leq b_i$ are stacked into one system; the Chebyshev centre of the intersection is found via the linear program

$$\max_{y, r} \quad \text{s.t.} \quad a_i^\top y + r \|a_i\|_2 \leq b_i \quad \forall i, \quad r \geq 0, \quad (5)$$

i.e. the centre of the largest ball that fits inside every half-space simultaneously. $r \leq 0$ means the per-year near-optimal spaces do not overlap at all — itself a meaningful finding (no design is simultaneously near-optimal in all years at that slack).

3 Implementation architecture

The feature lives in `scripts/near_optimal.py` (core algorithms, ~670 lines), `scripts/snakemake_near_optimal` (orchestration wrapper), the `explore_near_optimal` Snakemake rule, and an `mga:` block in `config.yaml`. It runs *after* a network is solved — it consumes the `*_OPT.nc` network and never alters the ordinary solve.

3.1 Dimension auto-derivation

A dimension is a technology; its value is the summed installed capacity (MW) of every *extendable*, *costed* (`capital_cost > 0`) component belonging to it. The set of selectable dimensions is **auto-derived** from the built network — nothing is hardcoded: flipping a technology to `expansion:`

`true` in GreenBubble’s own config automatically makes it available. The grouping key per component is GreenBubble’s own `comp_tech_map` (component name $\rightarrow$ technology), falling back to the PyPSA `carrier` attribute. The resulting structure is exactly the nested `weights` dictionary PyPSA’s own MGA API expects:

```
{tech_key: {"Generator": {"p_nom": {"comp_name": 1.0, ...}},
            "Store": {"e_nom": {"comp_name": 1.0, ...}}, ...}}
```

Users then pick a small subset (typically 4–5) by name; an unknown name raises an error listing what is actually available for that network.

3.2 The constraint-correctness problem

This is the central engineering issue we had to solve, and the reason the per-direction solve is *not* a direct call to PyPSA’s built-in MGA helpers. GreenBubble adds project-specific custom Linopy constraints to every ordinary solve (most importantly a maximum renewable-electricity-to-grid-sales constraint, `add_max_RE_sales_constraint`) via an `extra_functionality` hook. PyPSA’s native `n.optimize.optimize_mga / optimize_mga_in_direction` build their *own* model internally and **never call `extra_functionality`** — so using them unmodified would silently drop our custom constraints and compute the near-optimal space against a *looser*, incorrect feasible region (higher volumes, wrong must-have/must-avoid flags).

We therefore reproduce PyPSA’s own per-direction MGA procedure manually, step for step, and insert our constraint re-application in the middle:

```
def _solve_mga_in_direction(n, direction, dimensions, c_opt, slack, ...)
:
    """Constraint-correct replacement for
    n.optimize.optimize_mga_in_direction."""
    m = n.optimize.create_model()                # <- PyPSA
    cost_objective = m.objective                 # captured before
    overwrite
    apply_custom_constraints(n, m, n_flags=n_flags, # <- GreenBubble
                            re_alpha=re_alpha)
    _add_budget_constraint(m, cost_objective,      # <- GreenBubble
                           (Eq. 3)
                           c_opt, slack)

    m.objective = -sum(
        float(direction[k]) *
        n.optimize.build_linexpr_from_weights(dimensions[k], model=m) #
        <- PyPSA
        for k in dimensions if direction.get(k, 0)
    )

    status, condition = n.optimize.solve_model(...) # <- PyPSA
    coords = n.optimize.project_solved(dimensions) if status == "ok"
    else None
    # <- PyPSA

    return status, condition, coords
```

Every Tier-1, Tier-2, and Tier-3 solve funnels through this one function, so the near-optimal space is always computed against *exactly* the same feasible region as an ordinary GreenBubble solve.

3.3 Stochastic-network support

GreenBubble also supports a PyPSA stochastic (multi-scenario) optimisation mode, orthogonal to anything in either cited paper. Since the shared first-stage investment variables (p_{nom} , e_{nom}) are the same across scenarios, Tiers 1–2 work *transparently* on a stochastic network once the scenario-broadcast (`scenario`, `name`) MultiIndex on the static component tables is collapsed back to plain component names (`is_stochastic`, `_flat_static`); the MGA solve itself needs no change since `n.objective` is already the expected cost. Tier 3 is deliberately **skipped** on a stochastic network: its scenarios already span multiple years/conditions, so intersecting *again* across years would double-count the same uncertainty axis.

4 What we reused versus what we wrote from scratch

This is the section most relevant to our discussion. We deliberately built as much as possible on top of **PyPSA’s own native MGA machinery** (which is Fabian Neumann’s reference implementation of Neumann & Brown 2021, shipped inside PyPSA itself) rather than re-deriving low-level LP mechanics. **We did not have access to a Grochowicz-authored code repository for the 2023 paper** — our reference for the robustness extension is the published paper’s equations and section numbers only. Below is an honest, function-level account of provenance.

4.1 Reused directly via PyPSA’s public API (not copied, called)

PyPSA function	Module	Used for
<code>n.optimize.create_model()</code>	<code>pypsa</code> <code>.optimization</code>	Building the Linopy LP model for each MGA solve
<code>n.optimize.build_linexpr_from_weights()</code>	<code>pypsa</code> <code>.optimization</code>	Turning a σ -weights dict into a Linopy linear expression — used for both the direction objective and Tier 3’s “fix aggregate capacity” constraint
<code>n.optimize.solve_model()</code>	<code>pypsa</code> <code>.optimization</code>	Solving each direction’s LP with the chosen solver/profile
<code>n.optimize.project_solved()</code>	<code>pypsa</code> <code>.optimization</code>	Reading solved capacities back, projected onto the selected dimensions
<code>generate_directions_haltan/evenly_spaced/random()</code>	<code>pypsa</code> <code>.optimization</code> <code>.mga</code>	Sampling search directions for Tier 2 (used unmodified via a thin wrapper)

The `weights` dictionary schema itself (Section 3.2) is also PyPSA’s own MGA API convention, adopted as-is so our dimension registry interoperates with these functions without translation.

4.2 Implemented from scratch (GreenBubble-original)

- **The manual per-direction solve** (`_solve_mga_in_direction`), reproducing PyPSA’s `optimize_mga_in_direction` step-by-step specifically to insert `apply_custom_constraints` — the fix described in Section 3.2. This is the one piece of PyPSA-adjacent logic we *did* re-derive, and only because the public API leaves no hook to extend it instead.
- **The sign-robust budget** (Eq. 3) generalising Neumann & Brown’s $(1 + \varepsilon) c_{\text{opt}}$ to a price-driven, potentially negative objective.
- **Dimension auto-derivation** (`available_dimensions`, `resolve_dimensions`) from GreenBubble’s own technology config and component-to-technology mapping — no equivalent exists in PyPSA or in the cited papers, which assume dimensions are given.
- **Stochastic-network transparency** (Section 3.3) — an extension neither paper addresses, since PyPSA’s scenario API postdates both.

- **Tier 3 in full:** the shared-bound anchor c^* (eq. 9), and the Chebyshev-centre LP itself (`scipy.optimize.linprog` over the stacked hull facet inequalities) — implemented from scratch in Python/NumPy/SciPy directly off the paper’s §2.3–2.4 equations, since no reference code was available to adapt.
- **Everything operational around the algorithm:** the Snakemake rule and wrapper, the three-tier `mga`: config block, CSV/JSON output schema, and all plotting (`plot_ranges`, `plot_hull_projections`, `plot_robustness` — see Section 6).

4.3 Verification against the reference implementation

While preparing this note we located the official code release for the 2023 paper: <https://github.com/aleks-g/intersecting-near-opt-spaces> (tag v1.0.1, GPL-3.0-or-later, Koen van Greevenbroek & Aleksander Grochowicz). We compared our implementation against it directly, function by function:

- **Chebyshev centre — confirmed identical.** Their `geometry.ch_centre()` solves precisely the same LP as our `chebyshev_centre()`: $\max r$ s.t. $a_i^\top x + r \|a_i\| \leq b_i$ over a hull’s `ConvexHull.equations`, same sign handling of the `qhull` convention. We arrived at this independently from the paper’s §2.4, and it matches their reference code exactly. Their version solves the LP with Gurobi (`gurobipy`) directly; ours uses `scipy.optimize.linprog`, a deliberate solver-agnostic choice consistent with GreenBubble’s documented no-Gurobi-licence fallback path (Section 6, validation notes) — same formulation, different backend.
- **Shared robustness bound — confirmed identical.** Their `calc_obj_bound.py` computes exactly `obj_bound = (1 + eps) * max(objs)`, i.e. our $c^* = \max_i c_{\text{opt}}(i)$ (eq. 9), letter for letter (up to our sign-robust $|c_{\text{opt}}|$ generalisation, which their code has no analogue for — their objective is always a positive investment cost, never GreenBubble’s negative price-mode net revenue).
- **Per-direction MGA solve — same algorithm, necessarily different API.** Their `utilities.solve_network_in_direction` does the same thing ours does (near-optimal budget as an `extra_functionality` constraint, custom direction objective, solve) but against `pypsa.linopf/linopt` — PyPSA’s pre-Linopy optimisation API. GreenBubble runs on PyPSA 1.0.7’s Linopy-based `n.optimize` API, which that code cannot run against unmodified; our `_solve_mga_in_direction` is the same algorithm ported to the current API, not a literal reuse.
- **Tier 2 (hull) — a real, disclosed difference.** Their `compute_near_opt.py` is substantially more sophisticated than ours: an *adaptive* direction sampler (facet normals of the current hull, directions touching the current Chebyshev ball, or both combined) with volume- or centre-shift convergence criteria, run until convergence or a direction budget is exhausted — versus our simpler fixed-count Halton/random/evenly-spaced sampling via PyPSA’s built-in generators. Theirs converges to a tighter hull for the same solve budget; ours is simpler to reason about and bounded by `n_directions`. See Section 4.4 — closing this gap is exactly what we are doing next.

4.4 Reuse plan: `near-optimal_dev3`

Their `geometry.py` module (Chebyshev centre, hull intersection, facet-normal/hypersphere-based direction sampling) is pure NumPy/SciPy/(Gurobi) geometry with **no dependency on the PyPSA optimisation API** — unlike the per-direction solve, it does not need porting to work with GreenBubble’s PyPSA version, only the LP backend needs adapting (Gurobi → solver-agnostic, as we already do for the Chebyshev centre). This makes it the practical, low-risk target for direct code reuse. We are vendoring it on a new `near-optimal_dev3` branch: their

geometry utilities are being carried over under their original GPL-3.0-or-later licence and attribution into a clearly separated module (rather than merged into GreenBubble’s MIT-licensed files, which GPL licensing does not permit), so that Tier 2’s direction sampling and Tier 3’s intersection can reuse their adaptive algorithm directly instead of our simpler from-scratch version — as much of their code reused as the licence boundary allows, properly cited.

4.5 The honest summary

We reused PyPSA’s low-level MGA *mechanics* as much as the public API allowed, and reimplemented the one piece it doesn’t expose a hook for (constraint-correct solving). Grochowicz et al. (2023)’s specific contribution — multi-year robustness via intersection and the Chebyshev centre — was implemented from the paper’s mathematics alone in the current released version, and Section 4.3 confirms it is mathematically identical to their reference code where the two are directly comparable. **We would value your read on Section 6’s Tier 2 comparison and the near-optimal_dev3 reuse plan above** — in particular whether the adaptive sampler’s convergence behaviour in `compute_near_opt.py` has any gotchas we should know about before we lean on it.

5 Validation status

- **Tier 1 (ranges):** validated on a synthetic demand-mode network, a real deterministic price-mode network ($c_{\text{opt}} \approx -60\text{M}$, i.e. net revenue), and a real 3-scenario stochastic price-mode network ($c_{\text{opt}} \approx -13.3\text{M}$ expected). In every case the optimum stayed feasible under its own budget, the custom RE-to-grid constraint was confirmed present in the MGA model (scenario-aware where relevant), and $\min \leq \text{optimal} \leq \max$ held for every dimension.
- **Tiers 2–3:** validated on synthetic networks (hull geometry, intersection/Chebyshev-centre correctness); Tier 2 additionally exercised end-to-end on a real network for the tutorial results shown in Section 6.
- **Solver sensitivity (an operational finding worth flagging):** the MGA objective places weight on only one or two capacity variables, so the LP is highly degenerate (many equally-optimal vertices). On a real GreenBubble network, `highs-fast` (barrier without crossover) was observed to return an outright *infeasible* status with a large primal–dual objective error — silently failing a direction rather than just under-reporting its extent. `gurobi-barrier-fast` did not show this failure across 100+ direction solves and remains fast, so it is our recommended default; without a Gurobi licence we fall back to a profile that terminates at a vertex (`highs-simplex` / `highs-default`) rather than an interior barrier point.

6 Example results (Tutorial 5)

GreenBubble’s tutorial 5 (`tutorials/5_near_optimal`) runs all three tiers end-to-end on a small demonstration network with dimensions `[onwind, solar, electrolysis, methanolisation, battery]` at `slack = 0.05`. The following figures are produced automatically by `scripts/near_optimal.py`’s plotting functions.




Figure 1: Tier 1 — per-technology near-optimal capacity band (blue) with the cost-optimal capacity marked (red). A technology whose band's lower bound sits above zero is must-have at this slack; one whose upper bound sits at zero is must-avoid.




Figure 2: Tier 2 — pairwise 2-D projections of the sampled extreme points and their convex hull, with the cost optimum marked (red star). Each point is one full model solve in a Halton-sampled direction.




Figure 3: Tier 3 — per-year near-optimal hulls overlaid, with the Chebyshev centre of their intersection (black X) as the design that stays near-optimal simultaneously across all years explored.

7 Discussion points for the meeting

1. We're vendoring `geometry.py`'s adaptive direction sampler (facets / touching-ball / combined, with volume- or centre-shift convergence) into GreenBubble on a `near-optimal_dev3` branch, to close the Tier 2 gap in Section 4.3 — any practical gotchas from running it at scale (angle-tolerance tuning, typical iteration counts, cases where it stalls) we should know before we lean on it?
2. Does the sign-robust budget generalisation (Section 2.2) for negative-objective (net-revenue, price-driven) systems come up in other applications you're aware of? We haven't seen it addressed explicitly in either cited paper or in the reference code (whose objective is always a positive investment cost).
3. Any feedback on the solver-degeneracy finding in Section 6 (barrier-without-crossover returning *infeasible* rather than just an interior point on a highly degenerate MGA objective) — has this come up in your own MGA runs? (The reference code always solves via Gurobi's default settings for `ch_centre/is_nonempty`, so may not hit this the same way.)
4. How would you extend Tier 3's per-year intersection to compose with PyPSA's native *stochastic* (probability-weighted scenario) mode, which we currently treat as orthogonal (Tier 3 skipped on stochastic networks, see Section 3.3) rather than combined?

References

- Neumann, F. & Brown, T. (2021). *The Near-Optimal Feasible Space of a Renewable Power System Model*. Electric Power Systems Research, 190, 106690. <https://doi.org/10.1016/j.epsr.2020.106690> (arXiv:1910.01891)
- Grochowicz, A., van Greevenbroek, K., Benth, F.E. & Zeyringer, M. (2023). *Intersecting Near-Optimal Spaces: European Power Systems with More Resilience to Weather Variability*. Energy Economics. <https://doi.org/10.1016/j.eneco.2022.106496> (arXiv:2206.12242)
- PyPSA MGA documentation: <https://docs.pypsa.org/latest/user-guide/optimization/modelling-to-generate-alternatives/>
- van Greevenbroek, K. & Grochowicz, A. Reference implementation for the 2023 paper. <https://github.com/aleks-g/intersecting-near-opt-spaces> (tag v1.0.1, GPL-3.0-or-later).
- GreenBubble docs: `docs/guide_near_optimal.rst`, `docs/tutorial_5_near_optimal.rst`, `scripts/near_optimal.py` (implementation), `scripts/snakemake_near_optimal.py` (orchestration).